

ECE 459: Programming for Performance

Assignment 3

Kenan Ali
Lawrence Park

March 27, 2013

Part 1. Baseline Performance

Majority of Runtime

Using perf on lynch (ece459-1), and on a personal machine (Core 2 T5250) we see that a majority of the runtime is spent in libm-2.16.so _ieee754_pow_sse2 (33.69%), in libz.so.1.2.7 (22.90%), and in test_harness Model::morph(int, double, double, double) (20.39%). The following is an excerpt from perf report:

Samples: 106K of event 'cycles', Event count (approx.): 97480443237			
44.31%	test_harness	libm-2.17.so	[.] 0x0000000000014c2d
22.90%	test_harness	libz.so.1.2.7	[.] 0x0000000000003766
20.39%	test_harness	test_harness	[.] Model::morph(int, double, double, double)
3.79%	test_harness	libQtGui.so.4.8.4	[.] QVector2D::length() const
2.81%	test_harness	libpng15.so.15.13.0	[.] 0x0000000000022525
1.34%	test_harness	libm-2.17.so	[.] pow

(Since lynch was not resolving the symbols, we used the personal machine to verify that libm-2.16.so was indeed calling _ieee754_pow_sse2.)

Since we were only allowed to modify model.cpp, we needed to look at its structure to determine where most of the time is spent within this function call. So we first ignored the library calls also dominating the runtime. Upon inspecting morph() with perf annotate, we see a lot of Qt-related function calls, particularly in the function's nested for-loops. These for-loops were determined to have the the critical loop with regards to performance, so we knew that our optimizations had to target this. The following excerpt shows a particularly hot portion of the critical loop within morph():

		QVector2D pQP(QP.y(), -QP.x());
		// Calculate u, v
		u = QVector2D::dotProduct(XP, QP) / QP.lengthSquared();
0.37		divsd %xmm0,%xmm1
12.98		movsd %xmm1,-0x78(%rbp)
		v = QVector2D::dotProduct(XP, pQP) / QP.length();
0.58		callq QVector2D::dotProduct(QVector2D const&,
		QVector2D const&@plt
		lea -0x60(%rbp),%rdi
		movsd %xmm0,-0x88(%rbp)
		callq QVector2D::length() const@plt
		movsd -0x88(%rbp),%xmm2

Finally, since the library function calls listed in perf report were significant chunks of the runtime, and because we were only allowed to modify morph(), we had to determine if morph() calls the other dominating components. By gaining this knowledge, we thought we might be able to address the amount of time being spent in these library functions, hence improving performance further.

Determine Library Callers

After examining the morph function in model.cpp, we saw a lot of pow() calls on double value types. Since the morph function is taking up a major part of the runtime, it can be reasonably assumed that the large amount of time spent in the pow calculations is due to this. We can verify this by looking through perf annotate on morph():

	inline const QPoint operator-(const QPoint &p1, const QPoint &p2)
	{ return QPoint(p1.xp-p2.xp, p1.ypp2.ypp); }
0.01	sub -0x94(%rbp),%r12d
0.44	movsd -0x78(%rbp),%xmm0
	sub -0x98(%rbp),%r13d
0.00	ucomis %xmm1,%xmm0
0.48	jbe 204
	movsd0xbdf(%rip),%xmm0
	# 402958 <_IO_stdin_used+0xa8>
0.00	ucomis -0x78(%rbp),%xmm0
0.36	ja 598
	else if(u <= 0) dist = sqrt(pow(X.x()
	- P.x(), 2.0) + pow(X.y() - P.y(), 2.0));
	204: xorpd %xmm3,%xmm3
0.00	ucomis -0x78(%rbp),%xmm3
0.28	jae 580
	else dist = sqrt(pow(X.x() - Q.x(), 2.0)
	+ pow(X.y() - Q.y(), 2.0));
	mov -0x98(%rbp),%eax
	sub %r15d,%eax
	cvtssi2 %eax,%xmm1
0.10	mov -0x94(%rbp),%eax
0.00	sub -0x9c(%rbp),%eax
	cvtssi2 %eax,%xmm0
0.13	230: mulsd %xmm0,%xmm0
	mulsd %xmm1,%xmm1
0.17	addsd %xmm0,%xmm1
0.08	sqrtsd %xmm1,%xmm0
0.12	ucomis %xmm0,%xmm0
0.17	jp 838
0.08	24a: movsd %xmm0,-0x78(%rbp)

To clarify the point: we see lots of XMM register use near pow() calls, so can be reasonably certain that a significant portion of the SSE-based double math originates from morph().

Unfortunately, because the libz symbols were not found, we are unable to find out exactly where these calls originate from. However, we can reason that they are not from morph() because morph carries out integer and double math not anything related to compression. It is likely that these calls are from the image generation stage of the algorithm, and are therefore not a concern for us.

Part 2. Improvements

1. Loop Parallelization via OpenMP

After inspecting the baseline profile, we identified two for-loops in the `morph()` function in which exist the critical loop. The critical loop is defined as embodying the program logic of the morphing algorithm, and performs floating-point computations on variables. The iterations from this loop operate on individual pixels, and do not mutate any global data. Parallelizing the execution of the loop via OpenMP involves minor synchronization overhead. By using OpenMP, our implementation parallelizes the critical loop by flattening the two loops (via the collapse clause). The following figure illustrates the changes made on top of the baseline implementation:

<i>Our Implementation</i>	<i>Baseline Implementation</i>
<pre>#pragma omp parallel for collapse(2) for(int i=0; i<wimg; ++i) { for(int j=0; j<himg; ++j) { QPoint X(i, j); double u, v; //double ww[lines]; //QPoint pp[lines]; QPoint sum(0.0, 0.0); double wsum = 0; // for each line for(int k=0; k<lines; ++k) { QPoint P = points[k]; QPoint P2 = points[k+lines]; //QPoint Q = points[k+2*lines]; //QPoint Q2 = points[k+3*lines];</pre>	<pre>for(int i=0; i<wimg; ++i) { for(int j=0; j<himg; ++j) { QPoint X(i, j); double u, v; double ww[lines]; QPoint pp[lines]; // for each line for(int k=0; k<lines; ++k) { // get original lines from reference line QPoint P = listLines[h]->at(k).first; QPoint Q = listLines[h]->at(k).second; QVector2D XP(X - P); QVector2D QP(Q - P);</pre>

Figure 1: Use of OpenMP work-sharing construct on the critical loop

To profile the loop parallelization, we simply calculated the mean of the total runtime of baseline implementation runs and the parallelized loop implementation. The following table illustrates the speedup gained exclusively from the loop parallelization:

Runtime	Baseline Code	Parallelized Code
Run 1	25.791	10.515
Run 2	25.539	10.422
Run 3	25.435	11.307
Run 4	25.468	12.006
Run 5	26.392	14.920
Run 6	27.402	12.898
Average:	26.004	12.011

Table 1. Profiled runtime of baseline code and parallelized code

The average speed-up from the loop parallelization was over 2.16. This significant speedup can be attributed to the fact that the critical loop of the morphing algorithm constitutes a majority of

the total runtime, and the algorithm is embarrassingly parallel due to the lack of data dependency between the loop iterations. As a result, parallelizing the critical loop incurs little synchronization overhead, and the workload to each thread of execution is reduced by the number of physical cores on the host system.

2. Pulling Redundant Code from the Critical Loop

As previously discussed in this report, the critical loop located inside the `morph()` function is responsible for the majority of the runtime. We noticed that the inner loop, which iterates over each line in the global vector variables, performs various computations that could be done prior to entering the critical loop. Specifically we separated the instantiation and accessing of various Qt data types from the math performed on them. The following figure illustrates the code changes made to the baseline code:

Optimized Implementation

```
void Model::morph(int h, double VARA, double VARB, double VARP) {
    //int n = 0;
    int lines = listLines[0]->size();

    QPoint points [lines*4];
    QVector2D vectors [lines*4];
    double lengths [lines*4];

    //pulled from the critical loop
    for(int k=0; k<lines; ++k) {
        QPoint P = listLines[h]->at(k).first;
        QPoint Q = listLines[h]->at(k).second;
        QPoint P2 = listAux[h]->at(k).first;
        QPoint Q2 = listAux[h]->at(k).second;

        QVector2D QP(Q - P);
        QVector2D pQP(QP.y(), -QP.x());
        QVector2D Q2P2(Q2 - P2);
        QVector2D pQ2P2(Q2P2.y(), -Q2P2.x());

        double QPx = QP.x();
        double QPy = QP.y();

        lengths[k] = QPx*QP.x + QPy*QP.y;
        lengths[k+lines] = QP.length();
        lengths[k+2*lines] = Q2P2.length();
        lengths[k+3*lines] = pow(QP.length(), VARA);

        points[k] = P;
        points[k+lines] = P2;
        points[k+2*lines] = Q;
        points[k+3*lines] = Q2;

        vectors[k] = QP;
        vectors[k+lines] = pQP;
        vectors[k+2*lines] = Q2P2;
        vectors[k+3*lines] = pQ2P2;
    }
}
```

Baseline Implementation

```
void Model::morph(int h, double VARA, double VARB, double VARP) {
    int n = 0;
    int lines = listLines[0]->size();

    for(int i=0; i<wing; ++i) {
        for(int j=0; j<himg; ++j) {

            QPoint X(i, j);
            double u, v;

            double ww[lines];
            QPoint pp[lines];

            // for each line
            for(int k=0; k<lines; ++k) {

                // get original lines from reference line
                QPoint P = listLines[h]->at(k).first;
                QPoint Q = listLines[h]->at(k).second;

                QVector2D XP(X - P);
                QVector2D QP(Q - P);

                QVector2D pQP(QP.y(), -QP.x());

                // Calculate u, v
                u = QVector2D::dotProduct(XP, QP) / QP.lengthSquared();
                v = QVector2D::dotProduct(XP, pQP) / QP.length();

                // get interpolating lines from reference line
                QPoint P2 = listAux[h]->at(k).first;
                QPoint Q2 = listAux[h]->at(k).second;

                QVector2D Q2P2(Q2 - P2);
                QVector2D pQ2P2(Q2P2.y(), -Q2P2.x());

                QVector2D X2 = QVector2D(P2) + u * Q2P2 + (v * pQ2P2) / Q2P2.length();

                QPoint p = X2.toPoint() - X;
            }
        }
    }
}
```

Figure 2: Use of OpenMP work-sharing construct on the critical loop

The code changes come from the realization that computations of inner loop variables, such as `QP`, `pQP`, `Q2P2`, and `pQ2P2`, do not depend on the iterations of the critical loop. This means that the baseline implementation repeatedly computes the same values for these variables in the critical loop. By pre-computing the inner loop variables in a separate loop and storing the values on arrays, the program runs faster without changing the semantics of the code. The following table tabulates the speedups gained by this change:

Runtime	Baseline Code	Parallelized Code
Run 1	25.791	17.958
Run 2	25.539	18.025
Run 3	25.435	18.245
Run 4	25.468	18.396
Run 5	26.392	18.546
Run 6	27.402	18.604
Average:	26.004	18.296

Table 2. Profiled runtime of baseline code and modified code

The speedup for the change was roughly 1.4. This modest speedup makes sense because the bulk of time-consuming computations, such as the QVector2D dot product and floating point divisions, are still taking place inside the critical loop. What was pulled from the critical loop constitutes simpler and unnecessarily repeated operations, such as constructing QPoint and QVector2D objects. The semantics of the program remain unchanged because the pre-computed variables are independent from the loop iterations on i and j.